

Name: Nour Aldeen Nabil

What is a Docker container, and how is it different from a virtual machine (VM)?

A container refers to a lightweight, stand-alone, executable package of a piece of software that contains all the libraries, configuration files, dependencies, and other necessary parts to operate the application. Since containers are isolated, they provide security, thus allowing multiple containers to run simultaneously on the given host

	Docker	Virtual Machines (VMs)
Boot-Time	Boots in a few seconds.	It takes a few minutes for VMs to boot.
Runs on	Dockers make use of the execution engine.	VMs make use of the hypervisor.
Memory Efficiency	No space is needed to virtualize, hence less memory.	Requires entire OS to be loaded before starting the surface, so less efficient.
Isolation	Prone to adversities as no provisions for isolation systems.	Interference possibility is minimum because of the efficient isolation mechanism.
Deployment	Deploying is easy as only a single image, containerised can be used across all platforms.	Deployment is comparatively lengthy as separate instances are responsible for execution.
Usage	Docker has a complex usage mechanism consisting of both third party and docker managed tools.	Tools are easy to use and simpler to work with.

What is the purpose of a Dockerfile? Explain the significance of directives like FROM, COPY, RUN, and CMD.

Dockerfiles are instructions. They contain all commands used to build an image. Each layer is a result of the changes from the previous layer.

- FROM: Initializes a new build stage and sets the Base Image
- RUN: Will execute any OS commands in a new layer to build the image
- CMD: Provides a default for an executing container. There can only be one CMD instruction in a Dockerfile
- COPY: Copies new files or directories from <src> and adds them to the filesystem of the container at the path <dest>.

Describe the layers of a Docker image. How does Docker optimize space and performance using these layers?

Docker images are built on a layered architecture, where each layer represents a filesystem delta. Layers are immutable and stack on top of each other to form the final image. The Dockerfile defines the instructions for building the image. Each instruction in the Dockerfile corresponds to a layer in the final image. During the image build process, Docker creates intermediate images at each step of the Dockerfile. These intermediate images encapsulate the changes made by each instruction.

Optimization Techniques:

Minimize Layers: To optimize build time and reduce image size, minimize the number of layers in the Dockerfile.

1. Combining multiple commands into a single RUN instruction reduces the number of intermediate images created.
2. Layer Caching: Docker caches intermediate layers to speed up subsequent builds. By structuring the Dockerfile intelligently and ordering commands from least likely to change to most likely to change, Docker can reuse cached layers efficiently.
3. Layer Reuse: When pulling images from a remote repository, Docker checks for existing layers in the local cache. Only missing layers are pulled from the remote repository, accelerating image downloads.

4. **Immutable Layers:** Embrace the immutable nature of Docker image layers. Once a layer is created, it cannot be modified or removed. Instead of modifying existing layers, build new layers on top to incorporate changes.

What are the benefits of using Docker volumes? Give an example where data persistence is crucial in a Docker container.

Docker volumes enable data persistence, ensuring data is retained even when containers are stopped or removed. They allow easy data sharing between containers, improve performance, separate application logic from data, and make backups straightforward. Volumes are also portable and managed efficiently by Docker, providing a reliable way to handle containerized data.

Example: Data Persistence in a Database Container

For example, in a PostgreSQL container, mounting a volume to the database's data directory ensures that data remains intact even if the container is stopped, updated, or replaced.

How does Docker handle networking? Explain the difference between bridge, host, and none network modes in Docker.

Docker handles networking by creating isolated networks where containers can communicate with each other and external systems. It uses network drivers to manage these connections.

1. The bridge network mode is the default, where each container gets its own IP address and communicates with other containers in the same bridge network via internal routing.
2. The host network mode removes network isolation and binds the container directly to the host machine's network stack, allowing the container to share the host's IP address and ports.
3. The none network mode disables networking entirely, isolating the container from any external or internal networks while still allowing loopback communication within the container.

Describe how you would configure container-to-container communication within a Docker network.

First we create a custom network using Docker's bridge driver, as it provides automatic DNS resolution for containers in the same network. When containers are started and attached to this network, they can communicate using their container names as hostnames.

Build the Docker image and verify its functionality:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
[root@server-01 multi_container]# curl 127.0.0.1:8000
<!DOCTYPE html>
<html>
<head>
  <meta charset='utf-8'>
  <meta http-equiv='X-UA-Compatible' content='IE=edge'>
  <title>Page Title</title>
  <meta name='viewport' content='width=device-width, initial-scale=1'>
  <link rel='stylesheet' type='text/css' media='screen' href='main.css'>
  <script src='main.js'></script>
</head>
<body>
  <h1>"Welcome to
    DevOps World!"</h1>
</body>
</html>
[root@server-01 multi_container]#
```

Ensure that the web server can communicate with the database container, and the database is persistent using a Docker volume:

```
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
[root@server-01 multi_container]# docker inspect web_app | grep IPAddress
"SecondaryIPAddresses": null,
  "IPAddress": "",
  "IPAddress": "10.0.0.2",
[root@server-01 multi_container]# docker inspect postgres | grep IPAddress
Error: No such object: postgres
[root@server-01 multi_container]# docker inspect postgres | grep IPAddress
postgres:14-alpine postgres_db
[root@server-01 multi_container]# docker inspect postgres_db | grep IPAddress
"SecondaryIPAddresses": null,
  "IPAddress": "",
  "IPAddress": "10.0.0.3",
[root@server-01 multi_container]#
```

```
[root@server-01 multi_container]# docker exec -it postgres_db bash
08c9cbcb4b4f:/# ping webapp
ping: bad address 'webapp'
08c9cbcb4b4f:/# ping web_app
PING web_app (10.0.0.2): 56 data bytes
64 bytes from 10.0.0.2: seq=0 ttl=64 time=0.063 ms
64 bytes from 10.0.0.2: seq=1 ttl=64 time=0.078 ms
64 bytes from 10.0.0.2: seq=2 ttl=64 time=0.052 ms
64 bytes from 10.0.0.2: seq=3 ttl=64 time=0.051 ms
64 bytes from 10.0.0.2: seq=4 ttl=64 time=0.054 ms
64 bytes from 10.0.0.2: seq=5 ttl=64 time=0.060 ms
```

Provide the command you used and explain how Docker resource limits can help in a production environment.

```
docker-compose.yml
• [root@server-01 multi_container]# docker inspect web_app | grep Memory
  "Memory": 536870912,
  "MemoryReservation": 0,
  "MemorySwap": 1073741824,
  "MemorySwappiness": null,
• [root@server-01 multi_container]# docker inspect web_app | grep Nano
  "NanoCpus": 1000000000,
```

Docker resource limits ensure that containers do not consume excessive resources, maintaining the stability and performance of the host system in a production environment. By setting CPU and memory constraints, you can prevent a single container from monopolizing resources, which is critical for environments with multiple containers running simultaneously. For instance, using options like `--memory` and `--cpus`, you can allocate specific limits to each container, ensuring fair resource distribution, protecting other containers and the host system from potential crashes or slowdowns caused by resource overuse, and optimizing overall performance and reliability.